



RISC-V Security Model (non-normative)

RISC-V Security Model Task Group

Version 0.1, 11/2023: This document is in development. Assume everything can change. See <http://riscv.org/spec-state> for details.

Table of Contents

Preamble.....	1
Copyright and license information.....	2
Contributors.....	3
1. Introduction.....	4
1.1. Requirements and tracking.....	4
1.2. Relationship to external profiles.....	4
2. RISC-V security model overview.....	6
2.1. Reference model.....	6
2.1.1. Trusted Computing Base (TCB).....	7
2.1.2. Root of trust.....	7
2.1.3. Isolation.....	8
2.1.4. Device assignment.....	9
2.1.5. Invasive subsystems.....	9
2.1.6. Event counters.....	9
2.1.7. Platform quality of service.....	9
2.1.8. Denial of service.....	10
2.2. Adversarial model.....	10
2.2.1. Logical.....	11
2.2.2. Physical and remote.....	13
2.3. Ecosystem security objectives.....	14
2.3.1. Secure identity.....	14
2.3.2. Security lifecycle.....	15
2.3.3. Attestable services.....	17
2.3.4. Authorized software.....	18
2.3.5. System updates.....	19
2.3.6. Isolation.....	21
2.3.7. Sealing.....	21
3. RISC-V security building blocks.....	23
3.1. Isolation.....	23
3.1.1. Privilege levels.....	23
3.1.2. Hypervisor extension.....	23
3.1.3. PMA.....	23
3.1.4. PMP and ePMP.....	24
3.1.5. sPMP.....	24
3.1.6. MMU.....	25
3.1.7. Supervisor domains.....	25
3.1.8. MTT.....	26
3.1.9. IOPMP.....	26

3.1.10. IOMTT	27
3.1.11. IOMMU	27
3.2. Software enforced memory tagging	28
3.3. Control flow integrity	28
3.4. Cryptography	28
3.5. Roadmap	29
3.5.1. Capability based architecture	29
3.5.2. Hardware enforced memory tagging	29
3.5.3. HFI	29
3.5.4. Lightweight isolation	29
3.5.5. System integration	29
3.5.6. Cryptography enhancements	29
4. Use case examples	30
4.1. Generic system without supervisor domains	30
4.1.1. Overview	30
4.1.2. Isolation model	31
4.1.3. Root of Trust	31
4.1.4. Authorized Boot	31
4.1.5. Attestation	31
4.1.6. Sealing	32
4.1.7. Device access control	32
4.2. Debug and performance management	32
4.3. Global Platform TEE	33
4.3.1. Overview	33
4.3.2. Isolation model	35
4.3.3. Root of Trust	36
4.3.4. Authorized boot	36
4.3.5. Attestation	37
4.3.6. Sealing	37
4.3.7. Device access control	37
4.3.8. System integration	38
4.4. Debug and performance management	38
4.5. Confidential computing on RISC-V (CoVE)	40
4.5.1. Overview	40
4.5.2. Isolation model	41
4.5.3. Root of trust	42
4.5.4. Authorized Boot	42
4.5.5. Attestation	42
4.5.6. Sealing	43
4.5.7. Device access control	43
4.5.8. System integration	44

4.5.9. Trusted device assignment	44
4.5.10. Debug and performance management	45
4.5.11. Platform QoS	46
5. Cryptography	47
Appendix A: References	48
Bibliography	49

Preamble



This document is in the [Development state](#)

Assume everything can change. This draft specification will change before being accepted as informative, so implementations made to this draft specification will likely not follow the future informative specification.

Copyright and license information

This specification is licensed under the Creative Commons Attribution 4.0 International License (CC-BY 4.0). The full license text is available at creativecommons.org/licenses/by/4.0/.

Copyright 2023 by RISC-V International.

Contributors

This RISC-V specification has been contributed to directly or indirectly by (in alphabetical order): Ali Zhang, Andy Dellow, Carl Shaw, Colin O’Flynn, Dean Liberty, Dong Du, Deepak Gupta, Colin O’Flynn, Guerney Hunt, Luis Fiolhais, Manuel Offenberg, Markku Juhani Saarinen, Munir Geden, Mark Hill, Nicholas Wood, Paul Elliott, Ravi Sahita, Samuel Ortiz, Steve Wallach, Suresh Sugumar, Terry Wang, Victor Lu, Ved Shanbhogue, Yann Loisel

Chapter 1. Introduction

This specification provides guidelines for how RISC-V security building blocks can be used to build secure systems for different use cases. It is aimed at developers of RISC-V technical specifications, as well as at designers of secure RISC-V systems.

A few example uses cases based on commonly used deployment models are provided. These use cases are not intended to be exhaustive, or to act as protection profiles. They are intended as templates to be used as general guidelines, which can be applied to a wide variety of use cases.

The examples may be extended over time as required. Protection profiles for more specific use cases are expected to be provided within relevant certification bodies, or as separate RISC-V specifications if required.

1.1. Requirements and tracking

Where this specification makes formal recommendations, those are captured as trackable requirements using the following format:

ID#	Requirement
CAT_NNN	The CAT is a category prefix that logically groups the requirements and is followed by 3 digits - NNN - assigning a numeric ID to the requirement. The requirements use the key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "NOT RECOMMENDED", "MAY", and "OPTIONAL" that are to be interpreted as described in RFC 2119 when, and only when, they appear in all capitals, as shown here. When these words are not capitalized, they have their normal English meanings.

A requirement or a group of requirements may be followed by non-normative text providing context or justification for the requirement. The non-normative text may also be used to reference sources that are the origin of the requirement.

Trackable requirements are intended for ease of reference across dependant specifications.

1.2. Relationship to external profiles

For the purpose of this specification, external profiles apply to existing ecosystems or segments, but do not generally mandate implementations or architectures. This specification does not aim to establish new profiles. Its main purpose is to provide guidelines for how RISC-V security building blocks can be used to build RISC-V products which can comply with existing profiles.

Some profiles cover some or all of:

- Security reference architectures and taxonomy
- Hardware and software security requirements
- Interfaces and programming models

- Protection profiles and certification programmes
- Reference firmware/software

Other profiles are focussed on processes and methodology.

Examples of external profiles include:

Profile	Description
Global Platforms (GP)	Trusted execution environments(TEE) and trusted firmware for mobile, connected clients, and IoT. Secure element (SE) for tamper resistant storage of and operations on cryptographic secrets. SESIP certification. globalplatform.org/
Platform Security Architecture (PSA)	Platform security requirements for connected devices. PSA Certified. www.psacertified.org/
Trusted computing group (TCG)	Trusted platform module (TPM) and Device identifier composition engine (DICE) for trusted platforms. TCG certification. trustedcomputinggroup.org/
Confidential computing consortium	Common principles and protocols for protecting data in use (confidential computing). confidentialcomputing.io/
NIST	Widely used US standards for security processes, protocols and algorithms. Examples for the purposes of this specification: NISTIR 8259 - IoT device cybersecurity capability SP800-207 - Zero Trust Architecture www.nist.gov/

This is not an exhaustive list, more examples can be found in the reference section of this specification.

Chapter 2. RISC-V security model overview

The aim of this chapter is to define common taxonomies and principles for secure systems as used in the rest of this and other RISC V specifications. It is divided into the following sections:

- Reference model
Defines a set of generic hardware and software subsystems used in examples and use cases to describe secure systems.
- Adversarial model
Defines common attack types on secure systems, and identifies RISC-V extensions which can aid mitigation.
- Ecosystem security objectives
Defines common security features and functional guidelines, used to deploy trustworthy devices in an ecosystem.

2.1. Reference model

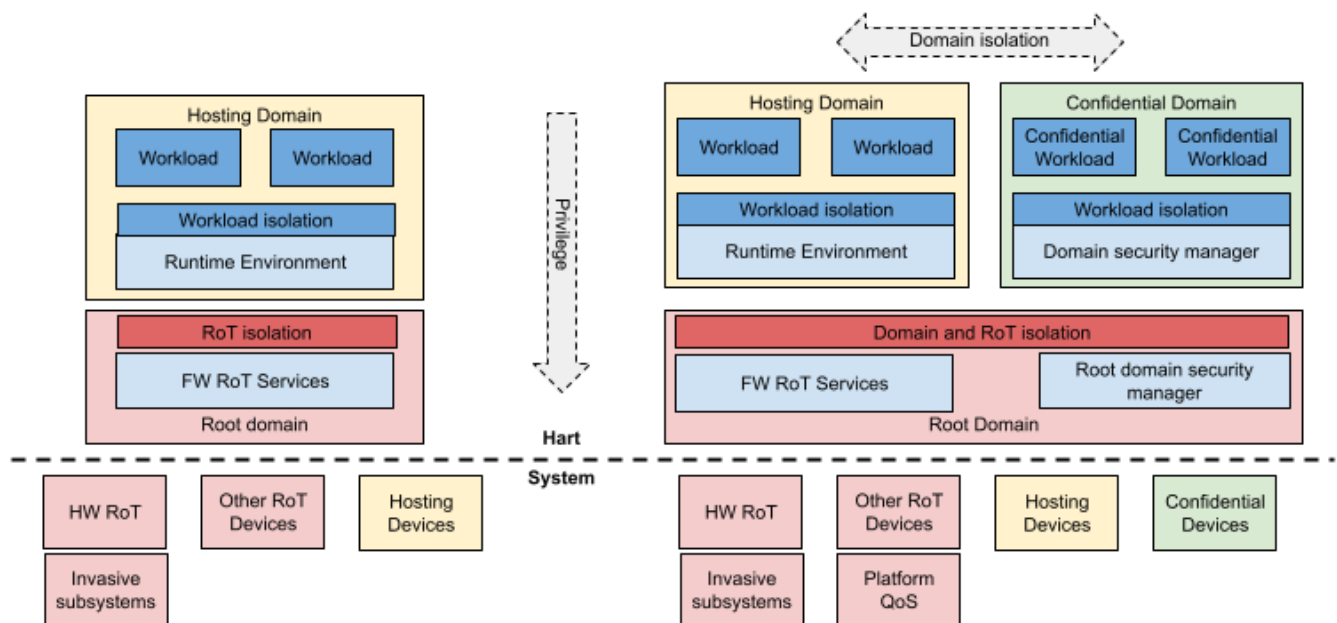


Figure 1: Generic security reference model

The figure above outlines a generic security reference model. It is not intended to describe any particular implementation and only aims to define a common taxonomy for the purpose of this and other RISC-V specifications.

Most systems are made up of different software components, often from different supply chains, each managing *assets* that need to be protected.

Examples of assets include:

- Cryptographic keys and credentials
- User data
- Proprietary models

- Secret algorithms

2.1.1. Trusted Computing Base (TCB)

The *Trusted Computing Base (TCB)* of a component is the minimum amount of software and hardware that needs to be trusted by that component.

The TCB of a component typically includes other software. For example:

- An operating system
- A hypervisor and a guest operating system
- A TEE security manager
- Hosting services such as orchestration and server provisioning software

2.1.2. Root of trust

The minimum amount of hardware and/or software that always has to be trusted on a system is its *root of trust (RoT)*. A RoT supports fundamental security services, for example:

- Boot and attestation
- Security lifecycle management
- Key derivations and sealing
- Security provisioning

Depending on use case and ecosystem requirements, a RoT can be:

- Hart firmware (FW RoT)
- A dedicated trusted subsystem (HW RoT), supporting a FW RoT

Using a HW RoT moves critical functions and assets off a Hart to a dedicated trusted subsystem, which can provide stronger protection against physical and logical attack than a complex Hart.



It is common for secure systems to support multiple trust chains with their own root of trust. For example, a TPM can be a root of trust for UEFI boot flows within a runtime environment, and a SIM can be a root of trust for user identity management.

For the purpose of this document, these should be treated as *secondary roots of trust*.

The HW RoT acts as a *primary root of trust* on the system. For example, the HW RoT governs boot, and can load firmware and manage the security lifecycle of a secondary root of trust.

ID#	Requirement
CAT_NNN	A complex secure system MUST implement a HW RoT

ID#	Requirement
CAT_NNN	A simpler system MAY not require a HW RoT but one is still recommended

Examples of a complex system include one with coherent multi-core, or out of order, application Harts.

Examples of a simpler system include a single core in-order microcontroller.



In this document, the terms "FW RoT" and "HW RoT" will be used as defined above. The term "RoT" on its own can be used where a rule or a rational applies to either model.

2.1.3. Isolation

Assets can be protected by *isolation*. Isolation reduces dependencies between components, and reduces the amount of software that needs to be trusted.

Isolation protects *resources*:

- Memory and memory mapped devices
- *Execution state*, including Hart register state

Examples of isolation mechanisms include:

- Privilege based isolation
More privileged software is able to enforce security guarantees for less privileged software.
- Physical memory isolation
More privileged software controls memory access for less privileged software.
- Domain isolation
Software in one domain cannot access or modify resources assigned to a different domain (without consent), regardless of privilege level.
(Higher privileged software in one domain cannot access resources assigned to a lower privilege level in a different domain)
- Virtualization
Virtualization creates and manages *virtual resources* - compute, memory, devices - independent of actual physical hardware. A system, or individual domains, can be virtualized.

On complex systems the TCB can grow large and get difficult to certify and attest.

Domain isolation enables confidential workloads to be separated from complex hosting software, including other workloads. The TCB of a confidential workload can be reduced to a domain security manager in a confidential domain, and the RoT, while allowing the main runtime environment in a separate hosting domain to remain in control of resource management.

Domain isolation use cases include:

- Platform security services - for example: secure storage, user identity management, payment clients, DRM clients

- Hosted confidential third party workloads

2.1.4. Device assignment

Isolation policy needs to extend to device assignment:

- Physical memory access control for device initiated transactions
- Virtual memory translation for virtualized device transactions
- Interrupt management across privilege and domain boundaries

These policies can be enforced by system level hardware, controlled by Hart firmware.

2.1.5. Invasive subsystems

Invasive subsystems include any system or Hart feature which could break security guarantees, either directly or indirectly. For example:

- External debug
- Power and timing management
- RAS (*reliability, accessibility, serviceability*)

ID#	Requirement
CAT_NNN	Invasive subsystems MUST be controlled, or moderated, by a RoT.

2.1.6. Event counters

Event counters are commonly used for performance management and resource allocation.

However, they can also pose a security risk. For example, one workload monitoring an operation in a different workload, or an operation by higher privilege software, could be able to reveal assets used in those operations.

ID#	Requirement
CAT_NNN	Lower privilege software MUST NOT be able to monitor higher privilege software.
CAT_NNN	Software in one domain MUST NOT be able to monitor software in a different domain, without consent.

2.1.7. Platform quality of service

Server platforms can provide *platform quality of service (QoS)* features, consisting of Hart and system hardware and firmware aimed at managing access to shared physical resources across workloads, minimizing contention. For example:

- Memory bandwidth management
- Cache allocation policies across workloads, including workload prioritization

- Hart allocation policies across workloads

These types of features rely on monitoring resource utilization of workloads, similar to event counters, and optimizing resource allocation policies.

ID#	Requirement
CAT_NNN	Lower privilege software MUST NOT be able to monitor higher privilege software.
CAT_NNN	Software in one domains MUST NOT be able to monitor software in a different domain, without consent.

2.1.8. Denial of service

The RISC-V security model is primarily concerned with protection of assets.

For example, a hosting environment is free to apply their own resource allocation policy to any workloads. Including denying service. This applies in the same way to confidential workloads.

ID#	Requirement
CAT_NNN	Lower privilege software MUST NOT be able to deny service to higher privilege software, or other isolated workloads at the same privilege level.

Higher privilege software must always be able to enforce its own resource management policy without interference. Including scheduling and resource assignment policies.

2.2. Adversarial model

For the purpose of this specification, the main goal of an adversary is to gain unauthorized access to *resources* - memory, memory mapped devices, and execution state. For example, to access sensitive assets, to gain privileges, or to affect the control flow of a victim.

In general, adversaries capable of mounting the following broad classes of attacks should be considered by system designers:

- Logical
The attacker and the victim are both processes on the same system.
- Physical
The victim is a process on a system, and the attacker has physical access to the same system. For example: probing, interposers, glitching, and disassembly.
- Remote
The victim is a process on a system, and the attacker does not have physical or logical access to the system. For example, radiation or power fluctuations, or protocol level attacks on connected services.

Attacks can be direct or indirect:

- Direct

An adversary gains direct access to a resource belonging to the victim. For example: direct access to a memory location or execution state, or direct control of the control flow of a victim.

- Indirect

An adversary can access or modify the content of a resource by a side channel. For example: by analyzing timing patterns of an operation by a victim to reveal information about data used in that operation, or launching row-hammer style memory attacks to affect the contents of memory owned by the victim.

- Chained

An adversary is able to chain together multiple direct and indirect attacks to achieve a goal. For example, use a software interface exploit to affect a call stack, and use that to take redirect the control flow of a victim.

This specification is primarily concerned with ISA level mitigations against logical attacks.

Physical or remote attacks in general need to be addressed at system, protocol or governance level, and may require additional non-ISA mitigations. However, some ISA level mitigations can also help provide some mitigation against physical or remote attacks and this is indicated in the tables below.

The required level of protection can vary depending on use case. For example, a HW RoT may have stronger requirements on physical resistance than other parts of an SoC.

Finally, this specification does not attempt to rate attacks by severity, or by adversary skill level. Ratings tend to depend on use case specific threat models and requirements.

2.2.1. Logical

ID#	Attack	Type	Description	Current RISC-V mitigations	Planned RISC-V mitigations
CAT_NN N	Unrestricted access	Direct Logical	Direct access to unauthorized resources in normal operation.	<ul style="list-style-type: none"> • RISC-V privilege levels • RISC-V isolation (for example: PMP/sPMP, MTT, supervisor domains) • RISC-V hardware virtualization (H extension, MMU) 	

ID#	Attack	Type	Description	Current RISC-V mitigations	Planned RISC-V mitigations
CAT_NNN	Transient execution attacks	Chained Logical	Attacks on speculative execution implementations.	<p>Known (documented) attacks, except Spectre v1, are specific to particular micro-architectures. Micro-architecture for RISC-V systems is implementation specific, but must not introduce such vulnerabilities.</p> <p>This is an evolving area of research.</p> <p>For example: Spectre and meltdown papers Intel security guidance Arm speculative vulnerability</p>	Fence.t, or similar future extensions, could at least partially mitigate against Spectre v1.
CAT_NNN	Interface abuse	Chained Logical	Abusing interfaces across privilege or isolation boundaries, for example to elevate privilege or to gain unauthorized access to resources.	<ul style="list-style-type: none"> • RISC V privilege levels • RISC-V isolation 	High assurance cryptography
CAT_NNN	Event counting	Direct Logical	For example, timing processes across privilege or isolation boundaries to derive information about confidential assets.	<ul style="list-style-type: none"> • Data-independent timing instructions • Performance counters restricted by privilege and isolation boundaries (sscofpmf, smcntrpmf) 	

ID#	Attack	Type	Description	Current RISC-V mitigations	Planned RISC-V mitigations
CAT_NN	Redirect control flow	Chained Logical	Unauthorized manipulation of call stacks and jump targets to redirect a control flow to code controlled by an attacker.	<ul style="list-style-type: none"> Shadow stacks (Zicfiss) Landing pads (Zicfilp) 	

2.2.2. Physical and remote

ID#	Attack	Type	Description	RISC-V recommendations
CAT_NN	Analysis of physical leakage	Direct or indirect Physical or remote	For example, observing radiation, power line patterns, or temperature.	<ul style="list-style-type: none"> Implement robust power management and radiation control Data Independent Execution Latency (Zkt, Zvkt)
CAT_NN	Physical memory manipulation	Direct Logical or physical	<ul style="list-style-type: none"> Row-hammer type software attacks to manipulate nearby memory cells Using NVDIMM, interposers, or physical probing to read, record, or replay physical memory Physical attacks on hardware shielded locations to extract hardware provisioned assets 	<ul style="list-style-type: none"> Implement robust memory error detection, cryptographic memory protection, or physical tamper resistance Supervisor domain ID, privilege level, or MTT attributes, could be used to derive memory encryption contexts at domain or workload granularity Provide a degree of tamper resistance

ID#	Attack	Type	Description	RISC-V recommendations
CAT_N NN	Boot attacks	Chained Logical or physical	<ul style="list-style-type: none"> • Glitching to bypass secure boot • Retrieving residual confidential memory after a system reset 	<ul style="list-style-type: none"> • Implement robust power management • Implement cryptographic memory protection with at least boot freshness • Code hardening (for example, see Riscure white paper)
CAT_N NN	Subverting supply chains	Remote	Infiltration or collusion to subvert security provisioning chains, software supply chains and signing processes, hardware supply chains, attestation processes, development processes (for example, unfused development hardware or debug authorizations)	Deploy appropriate governance, accreditation, and certification processes for an ecosystem.

2.3. Ecosystem security objectives

Ecosystem security objectives identify a set of common features and mechanisms that can be used to enforce and establish trust in an ecosystem.

These features are defined here at a functional level only. Technical requirements are typically use case specific and defined by external certification programmes.

In some cases RISC-V non-ISA specifications can provide guidance or protocols. This is discussed more in use case examples later in this specification.

2.3.1. Secure identity

ID#	Requirement
CAT_NNN	A security platform MUST be securely identifiable

Identifies the immutable part of the security platform - immutable hardware, configurations, and firmware. Immutable components cannot change after completed security provisioning (see also security lifecycle management).

A *secure identity* is one capable of generating a cryptographic signature which can be verified by a remote party. Usually an asymmetric key pair, but sometimes symmetric signing schemes can be

used). It is typically used as part of an attestation process.

Its scope and uniqueness depends on use case. For example:

- Unique to a system
- Shared among multiple systems with the same immutable security properties (group based anonymization)
- Anonymized using an attestation protocol supporting a third party anonymization service

It can be directly hardware provisioned, or derived from other hardware provisioned assets.

2.3.2. Security lifecycle

ID#	Requirement
CAT_NNN	A secure system MUST manage a security lifecycle.

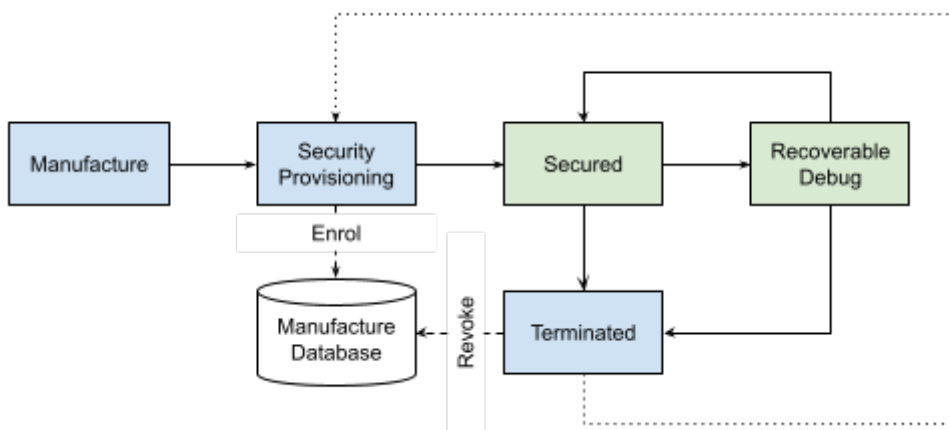


Figure 2: Generic security lifecycle

A security lifecycle reflects the trustworthiness of a system during its lifetime and reflects the lifecycle state of hardware provisioned assets.

It can be extended as indicated below to cover additional security provisioning steps such as device onboarding, device activation, user management, and RMA processes. These are use case or ecosystem specific and out of scope of this specification.

For the purpose of this specification, *revealing debug* includes any HW or FW debug capability which

- Could break security guarantees or could expose assets
- Is not part of an attested trust contract with a reliant party

Examples of revealing debug include revealing logging, external debug or boundary scans, dedicated debug builds of software components, or enabling self-hosted debug for a component.

Depending on use case, an attested software component can include debug capabilities managed through an ecosystem defined governance process - *trusted debug*. For example, self-hosted debug enabled following an ecosystem specific authorization process. In this case the debug capability,

and the associated governance, is part of the trust contract with a reliant party.

For the purpose of this specification, a minimum security lifecycle includes at least the following states:

- **Manufacture** - The system may not yet be locked down and has no hardware provisioned assets
- **Security provisioning** - The process of provisioning hardware provisioned assets
Depending on ecosystem requirement, security provisioning could be performed in multiple stages through a supply chain and may require additional sub-states. These types of application specific extensions are out of scope of this specification.
- **Secured** - hardware provisioned assets are locked (immutable), only authorized software can be used, and revealing debug is not enabled.
Additional specific provisioning stages can take place in this state - for example network onboarding and device activation, TSS/App/Device attestation or user identity management. This is out of scope of this specification.
- **Recoverable debug** - part of the system is in a revealing debug state
At least the RoT is not compromised and hardware provisioned secrets remain protected. This state is both attestable and recoverable. For example, revealing debug is enabled for a domain without compromising another domain or any RoT services.
- **Terminated** - any system change which could expose hardware provisioned assets

Typically hardware provisioned assets are made permanently inaccessible and revoked before entering this state. This also protects any derived assets such as attestation and sealing keys.

A system could support re-provisioning from a terminated state, for example following repair/RMA. This can be viewed as equivalent to starting over from the security provisioning state, and creates a new instance with a new secure identifier.

ID#	Requirement
CAT_NNN	Hardware provisioned assets MUST only be accessible while the system is in secured state, or a recoverable debug state.
CAT_NNN	Derived assets MUST only be available if a component is in secured state.

ID#	Requirement
CAT_NNN	Hardware provisioned assets MUST only be accessible while the system is in secured state, or a recoverable debug state (with the recoverable debug state in attestation evidence).
CAT_NNN	Derived assets MUST only be available if a component is in secured state.

A derived asset in this context is any asset derived from hardware provisioned assets. For example attestation keys, or sealing keys for a supervisor domain.

ID#	Requirement
CAT_NNN	Revealing debug MUST be reflected in attestation.

Attestable states are ones where the RoT and hardware provisioned assets are not compromised by debug and a valid attestation can be generated reflecting that state:

- Secured
- Recoverable debug

In other states the system is not able to generate a valid attestation key. It is still *indirectly attestable* as any generated attestation will not be signed correctly and can be rejected by a reliant party.

Trusted debug is part of a trust contract with a reliant party, and application specific. The presence of trusted debug can be determined indirectly by a reliant party through other attested properties, for example measurements.

2.3.3. Attestable services

For the purpose of this specification a confidential service can be any isolated component on a system. For example, a hosted confidential workload, or an isolated application security service.

ID#	Requirement
CAT_NNN	A confidential service, and all software and hardware components it depends on, MUST be attestable.

Attestation allows a remote reliant party to determine the trustworthiness of a confidential service before submitting assets to it.

- Verify the security state of a confidential service
- Verify the security state of all software and hardware a confidential service depends on
- Establish an attested secure connection to a confidential service

Attestation can be direct or layered.

- Direct
The whole system can be defined by a single security platform attestation. For example, can be used in vertically integrated connected IoT devices and edge devices.
- Layered
Enables parts of the attestation process to be delegated to lower privileged components.

Direct and layered attestation are discussed in more detail in use case examples later in this specification.

ID#	Requirement
CAT_NNN	A security platform attestation MUST be signed by a HW RoT, if present, or by a FW RoT
CAT_NNN	A security platform attestation MUST be signed using a hardware provisioned (directly or derived) secure identity

ID#	Requirement
CAT_NNN	A layered attestation MAY be signed by lower privileged software, itself attested by a security platform attestation
CAT_NNN	Layered attestations MUST be cryptographically bound such that a relying party can determine that they <ul style="list-style-type: none"> • Were generated on the same system • Are fresh.



Software interfaces should only support either direct attestation or layered attestation workflows, never both, to prevent impersonation attacks.

2.3.4. Authorized software

Running unauthorized software can compromise the security state of the system.

ID#	Requirement
CAT_NNN	A system in secured or recoverable debug states MUST only load authorized software.
CAT_NNN	A system in security provisioning state SHOULD only load authorized software.

Two complementary processes can be used to authorize software:

- Measuring
A measurement is a cryptographic fingerprint, such as a running hash of memory contents and launch state.
- Verification
Verification is a process of establishing that a measurement is correct (expected)

A boot process is typically layered, allowing software to be measured and verified in stages. Different measurement and verification policies can be employed at different stages. This is discussed further in use case examples later in this specification. The properties discussed below still apply to each stage.



Measurements can be calculated at boot (*boot state*), and sometimes also dynamically at runtime (*runtime state*). Measuring runtime state can be used as a robustness feature to mitigate against unauthorized runtime changes of static code segments. It is out of scope of this specification, though the principles discussed below can still be applied.

Verification can be:

- Local
A measurement is verified locally on the device.

- Remote

A measurement is verified by a remote provisioning service, or a remote reliant party.

Verification can be:

- Direct

The measurement is directly compared with an expected measurement from a signed authorization.

- Indirect

The measurement is included in derivations of other assets, for example sealing keys, binding assets to a measured state.

ID#	Requirement
CAT_NNN	A security platform MUST be measured.
CAT_NNN	A security platform MUST be verified, either directly or indirectly, before launching services which depend on the security platform.

Verification ensures the system has loaded authorized software

ID#	Requirement
CAT_NNN	A system MUST only use authorizations from trusted authority.

- Direct verification requires a signed image authorization from a trusted authority before loading an image
For example, a signed image, or a separately signed authorization message.
- Indirect verification requires a signed authorization from a trusted authority for migrating assets bound to a previously measured state
For example, a signed provisioning message.

Either way, only authorizations from trusted authorities should be used. For example, from a list of hardware provisioned or securely discovered trusted authorities.

ID#	Requirement
CAT_NNN	Local verification MUST be rooted in immutable boot code.

For example, ROM or locked flash, or rooted in a HW RoT itself rooted in immutable boot code.

2.3.5. System updates

Over time, any non-immutable component may need updates to address vulnerabilities or functionality improvements. A system update can concern software, firmware, microcode, or any other updatable component on a system.

ID#	Requirement
CAT_NNN	All components on a system which are not immutable MUST be updatable.

Immutable components include at least immutable boot code. Some trusted subsystems can also include immutable software to meet specific security certification requirements.

System updates are typically layered so that updates can target only parts of a system and not a whole system. The properties discussed below still apply to any system update.

ID#	Requirement
CAT_NNN	A system update MUST be measured and verified before launch.

See [Section 2.3.4](#).

A system update can be:

- Deferred
The update can only be effected after a restart of at least the affected component, and all of its dependents.
- Live
The update can be effected without restarting any dependent components.

ID#	Requirement
CAT_NNN	Updates affecting a security platform SHOULD be deferred.
CAT_NNN	Updates MAY be live if live update capability, and suitable governance, is part of an already attested trust contract between a reliant party and the system.

A system update changes the attested security state of the affected component(s), as well as that of all other components that depend on it. It can affect whether a dependent confidential service is still considered trustworthy or not, as well as affect any derived assets such as sealing keys.

ID#	Requirement
CAT_NNN	System updates MUST be monotonic
CAT_NNN	System updates SHOULD be robust against update failures

Earlier versions may be carrying known vulnerabilities, or may affect the safe operation of a system in other ways.

For example, using derived anti-rollback counters (counter tree) rooted in a hardware monotonic counter.

A system can still support recovery mechanisms, with suitable governance, in the case of update failures. For example, a fallback process or a dedicated recovery loader.

Success criteria for a system update are typically use case or ecosystem specific and out of scope of this specification. Examples include local watchdog or checkpoints, and network control through a secure update protocol, and a dedicated recovery loader.

ID#	Requirement
CAT_NNN	System updates, and authorization messages, SHOULD only be received from trusted sources.

A system update is itself always verified before being launched. Verifying the source as well can mitigate against attempts to inject adversary controlled data into a local update process. Including into protected memory regions.

2.3.6. Isolation

Complex systems include software components from different supply chains, and complex integration chains with different roles and actors. These supply chains and integration actors often share mutual distrust:

- Developed, certified, deployed and attested independently
- Protected from errors in, or abuse from, other components
- Protected from debugging of other components
- Contain assets which should not be available to other components

Use cases later in this specification provide examples of RISC-V isolation models.

ID#	Requirement
CAT_NNN	Isolated software components SHOULD be supported

An isolated component has private memory and private execution contexts not accessible to other components.

ID#	Requirement
CAT_NNN	Devices MUST not access memory belonging to an isolated component without permission

Isolation can also extend to other features, such as interrupts and debug.

2.3.7. Sealing

Sealing is the process of protecting confidential assets on a system.

ID#	Requirement
CAT_NNN	Sealed assets SHOULD only be possible to unseal in a secured state

Sealing can be:

- Local
Local sealing binds assets to a local device (hardware unique sealing).
- Remote
Remote sealing binds assets to credentials provided by a remote provisioning service following

successful attestation.

Local sealing can be:

- Direct
Direct sealing binds assets to sealing keys derived by trusted security services, or a HW RoT.
- Layered
Layered sealing enables delegation of some sealing key derivations to lower privileged software.

ID#	Requirement
CAT_NNN	Local sealed assets MUST only be possible to unseal on the same physical instance of a system that they were sealed on.

For example, using sealing keys derived from a hardware provisioned *hardware unique key (HUK)*.

Sealing is discussed further in use cases examples later in this document.

Chapter 3. RISC-V security building blocks

This chapter outlines brief descriptions of RISC-V security building blocks discussed in this specification, together with general guidelines and links to technical specifications.

See also the reference use cases chapter of this specification for common examples of how RISC-V security building blocks can be combined.

3.1. Isolation

Isolation enable separation of software components executing on a Hart, as well as device assignment. RISC-V enables:

- Privilege based isolation
- Physical memory access control (hart and device-initiated accesses)
- Virtual memory management (hart and device virtualization)
- Hypervisor extension
- Supervisor domains

3.1.1. Privilege levels

[Privileged ISA](#)

Standard privilege levels - Machine mode (M), Supervisor mode (S), and User mode (U) - enable separation of more privileged software from less privileged software.

3.1.2. Hypervisor extension

[Privileged ISA](#)

Hypervisor extension supports standard supervisor level hypervisors. It extends S mode into Hypervisor-extended supervisor mode (HS), and a virtual supervisor mode (VS) for guests. It also extends U mode into standard user mode (U) and virtual user mode (VU).

Isolation of guests is enforced using two-stage address translation and protection. Two-stage address translation and protection is in effect in VS and VU modes.

Alternatively sPMP can be used instead of MMU to support static partition hypervisors, for example on systems with hard and deterministic real time requirements [Note -The sPMP for Hypervisor extension has not been specified yet].

MMU, PMP/ePMP, and sPMP are discussed later in this chapter.

3.1.3. PMA

[Privileged ISA](#)

Physical memory attributes (PMA) are intended to capture inherent properties of the underlying hardware. For example, read-only ROM regions, or non-cachable device regions. Often PMA can be fixed at design time or at boot, but sometimes runtime PMA can be required.

A separate hardware checker - *PMA checker* - enforces PMA rules at runtime once a physical address is known. PMA rules are always checked on every physical access, and typically configured by region.

Other RISC-V mechanisms can also be used for managing memory attributes. For example [MTT](#) can provide dynamic meta-data during Sv address translation (MTT Svpam). Or [MMU TLB](#) can be used by a kernel or hypervisor to manage memory attributes and permissions at a virtualized system level.

Which mechanism is used depends on use case. They can be combined, in which case any memory access must pass all memory attribute checks.

3.1.4. PMP and ePMP

[Privileged ISA](#)

Physical memory protection (PMP) enables M-mode to access-control physical memory for supervisor or HS modes.

ID#	Requirement
CAT_NNN	PMP configurations MUST only be directly accessible to machine mode

Individual access controlled regions can be locked until the next system reset to create temporal isolation boundaries, such as protecting immutable boot code.

ePMP extends PMP protection by allowing machine mode to restrict its own access to memory allocated to lower privilege levels. This can be used to mitigate against privilege escalation attacks, for example.

ID#	Requirement
CAT_NNN	If PMP is supported then ePMP MUST be supported.

3.1.5. sPMP

github.com/riscv/riscv-smp

Supervisor PMP (sPMP) enables supervisor mode to control physical memory access for U mode.

sPMP allows supervisor mode to restrict its own access to memory allocated to lower privilege levels. This can be used to mitigate against privilege escalation attacks, for example.

When combined with H-extension, sPMP can be nested so that the hypervisor can control memory allocations to its guests, and each guest can control its own memory allocations to its workloads.

ID#	Requirement
CAT_NNN	Either sPMP, 1st-stage or G-stage page tables MUST be used to protect Supervisor domain H/S-mode from lower privilege levels.

3.1.6. MMU

Privileged ISA

Memory management unit (MMU) enables address translation and protection for:

- Isolating an OS from workloads on a system without H-extension (one-stage translation)
- Isolating a hypervisor from a guest, on a system with H-extension (two-stage translation)

ID#	Requirement
CAT_NNN	Either PMP/ePMP or MTT MUST be used to protect M-mode from lower privilege levels.

3.1.7. Supervisor domains

github.com/riscv/riscv-smmtt

Supervisor domains allow software components on the same hart to be developed, certified, deployed and attested independently of each other.

A supervisor domain is a compartment above M-mode, physically isolated - memory, execution state, and devices - from other supervisor domains regardless of privilege level (below M-mode). Isolation and context switching between supervisor domains are managed by M-mode firmware.

A supervisor domain is identified at architecture level by a *supervisor domain id (SDID)*, managed by M-mode firmware.

ID#	Requirement
CAT_NN	PMP/ePMP or MTT MUST be used to enforce physical memory isolation boundaries for supervisor domains, and to protect machine mode from any supervisor domain.

PMP can be used for more static and deterministic use cases.

MTT can be used where more fine grained dynamic resource management across supervisor domain boundaries is required.



MTT can be sufficient for protecting Root domain in the sense that M-mode can enforce that its own resources are never assigned to another domain. PMP/ePMP still add further protections for M-mode, such as the ability to implement temporal isolation boundaries within M-mode (for example, protect early boot code), or to prevent itself from accessing or executing from memory assigned to lower privilege levels (privilege escalation).

ID#	Requirement
CAT_NNN	A system supporting supervisor domains MUST support supervisor domain extensions for interrupts (Smsdia) and performance counters (TBD), and SHOULD support supervisor domain extensions for external debug (Smsdsd TBD).

Interrupts: github.com/riscv/riscv-smmtt

External debug: github.com/riscv-non-isa/riscv-external-debug-security

Performance counters:

These extensions enable management of interrupts, external debug, and performance counters across supervisor domain boundaries. M-mode firmware should context switch hart HPM event/counters to manage isolation of performance counters:

- External debug can be enabled for one supervisor domain without affecting other supervisor domains
- M-mode firmware manage interrupt routing and preemption across supervisor domain boundaries
- M-mode firmware can ensure that performance counters cannot be used by software in one supervisor domain to measure operations in other supervisor domains

3.1.8. MTT

github.com/riscv/riscv-smmtt

The *memory tracking table (MTT)* is a memory structure managed by machine mode, tracking memory ownership across supervisor domains. It is designed to enable fine grained dynamic memory management across supervisor domain boundaries, with policy typically set by a hypervisor in a hosting domain responsible for resource management.

ID#	Requirement
CAT_NNN	Either PMP/ePMP or MTT MUST be used to protect M-mode from lower privilege levels
CAT_NNN	MTT configurations MUST only be directly accessible to machine mode



MTT can be sufficient for protecting Root domain in the sense that M-mode can enforce that its own resources are never assigned to another domain. PMP/ePMP still add further protections for M-mode, such as the ability to implement temporal isolation boundaries within M-mode (for example, protect early boot code), or to prevent itself from accessing or executing from memory assigned to lower privilege levels (privilege escalation).

3.1.9. IOPMP

github.com/riscv-non-isa/iopmp-spec

IOPMP is a system level component providing physical memory access control for device-initiated transactions, complementing PMP and sPMP rules.

ID#	Requirement
CAT_NNN	A system which supports PMP/ePMP, or sPMP, MUST implement IOPMP for device access control.
CAT_NNN	IOPMP configurations MUST only be directly accessible to machine mode.



IOPMP defines multiple "models" for different system configurations. Unless specified differently in the use cases in this specification, system designers are free to choose any IOPMP model.

3.1.10. IOMTT

github.com/riscv/riscv-smmtt

IOMTT is a system level component providing physical memory access control for device-initiated transactions, complementing MTT rules.

ID#	Requirement
CAT_NNN	A system which supports MTT MUST implement IOMTT for access-control for device-initiated memory accesses.
CAT_NNN	IOMTT configurations MUST only be directly accessible to machine mode.
CAT_NNN	A system which implements IOMTT MAY also implement IOPMP to access-control device-initiated access to M-mode memory.



IOMTT can also be sufficient for protecting Root devices in the sense that M-mode can enforce that its own resources are never assigned to another domain. Use of IOPMP or similar still adds further protections. For example, a system may require that Root devices are not able to access memory assigned to TEE domain.

3.1.11. IOMMU

github.com/riscv-non-isa/riscv-iommu

IOMMU is a system level component performing memory address translation from IO Virtual Address to Physical Address, allowing devices to access virtual memory locations. It complements MMU configurations.

ID#	Requirement
CAT_NNN	Systems supporting MMU SHOULD also support IOMMU
CAT_NNN	Systems supporting IOMMU MUST also enforce physical memory access control for M-mode memory against device-initiated transactions (IOMTT or IOPMP).

3.2. Software enforced memory tagging

github.com/riscv/riscv-j-extension

Memory tagging (MT), is a technique which can improve the memory safety of an application. A part of the effective address of a pointer can be masked off, and used as a tag indicating intended ownership or state of a pointer. The tag can be used to track accesses across different regions, as well as protecting against pointer misuse such as "use after free". The pointer masking should use the proposed J-extension pointer masking extension (Smpm, Smnpm, Ssnpm).

With software based memory tagging the access rules encoded in tags are enforced by software (compiler).

See also hardware enforced memory tagging below.

3.3. Control flow integrity

github.com/riscv/riscv-cfi

Control-flow Integrity (CFI) capabilities help defend against Return-Oriented Programming (ROP) and Call/Jump-Oriented Programming (COP/JOP) style control-flow subversion attacks, where an attacker attempts to modify return addresses or call/jump address to redirect a victim to code reused by the attacker.

These attack methodologies use code sequences in authorized modules, with at least one instruction in the sequence being a control transfer instruction that depends on attacker-controlled data either in the return stack or in memory used to obtain the target address for a call or jump. Attackers stitch these sequences together by diverting the control flow instructions (e.g., JALR, C.JR, C.JALR), from their original target address to a new target via modification in the return stack or in the memory used to obtain the jump/call target address.

RISC-V provides two defenses:

- Shadow stacks (Zicfiss) - protect return addresses on call stacks
- Labeled Landing pads (Zicfilp) - protect target addresses in jumps and branches

3.4. Cryptography

github.com/riscv/riscv-crypto

RISC-V includes ISA extensions in the following cryptographic areas:

- Scalar cryptography
- Vector cryptography
- Entropy source (scalar)

RISC-V cryptographic extensions are aimed at supporting efficient acceleration of cryptographic operations at ISA level. This can both help reduce the TCB of an isolated component, and avoid

hardware bottlenecks (for example, system level cryptographic subsystems).

The entropy source extension provides an ISA level interface to a hardware entropy source. Entropy source requirements can depend on use case or ecosystem specific requirements and RISC-V does not provide any entropy source technical specification. But the entropy source ISA specification does contain general recommendations and references.

ID#	Requirement
CAT_NNN	RISC-V systems SHOULD support either scalar or vector cryptographic ISA extensions
CAT_NNN	The entropy source ISA extension MUST be supported if either scalar or vector cryptographic ISA extensions are supported.

It is not necessary to support both scalar and vector operations, as a scalar operation can be viewed as a vector of size 1.

3.5. Roadmap

3.5.1. Capability based architecture

- Cheri
- Capstone

3.5.2. Hardware enforced memory tagging

Hardware enforced memory tagging extends software based memory tagging (see above) by moving enforcement of tagged rules to hardware.

3.5.3. HFI

Hardware-assisted fault isolation (HFI) aims to provide lightweight in-process isolation to mitigate against errors in one process compromising other processes within the same workload.

3.5.4. Lightweight isolation

- TBD

3.5.5. System integration

- WorldGuard

3.5.6. Cryptography enhancements

- PQ
- High assurance computing (HAC)

Chapter 4. Use case examples

This chapter provides a selection of example uses cases based on commonly used deployment models. These use cases are not intended to be exhaustive, or to act as protection profiles. They are intended as templates which can be used as general guidelines, which can be applied to a wide variety of use cases.

The examples may be extended over time as required. Protection profiles for more specific use cases are expected to be provided within relevant certification bodies, or as separate RISC-V specifications if required.

4.1. Generic system without supervisor domains

4.1.1. Overview

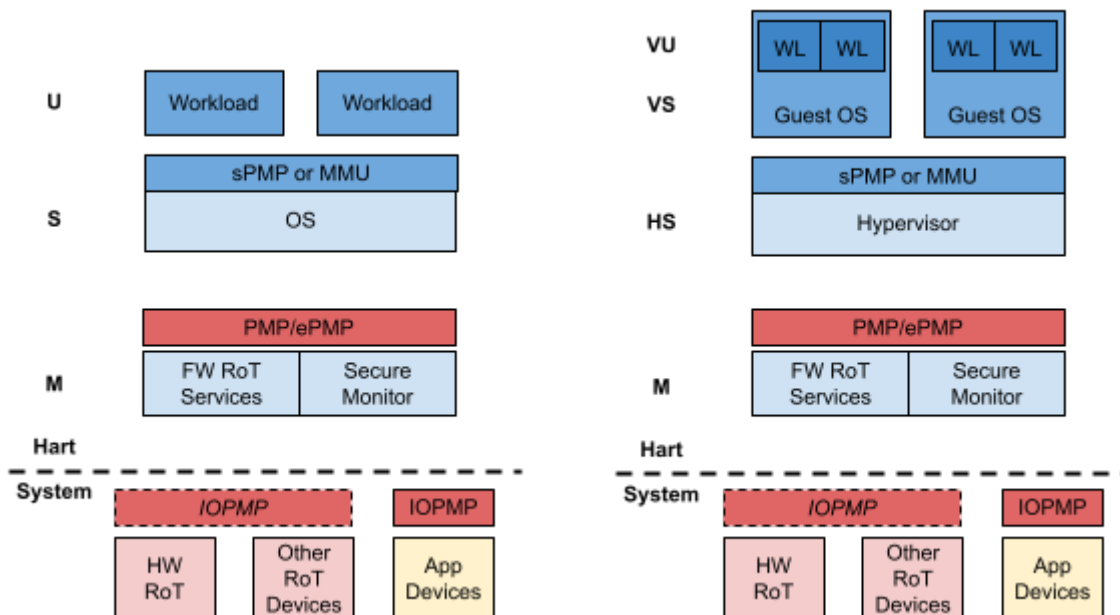


Figure 3: Generic vertically integrated system

A generic vertically integrated system can be either virtualized or non-virtualized.

M-mode hosts a FW RoT. An OS or a Hypervisor in S or HS mode controls applications or guests. Guests and applications execute in U or VS/VU modes and trust the OS or Hypervisor to provide isolation guarantees.

A system level HW RoT is recommended.



The RISC-V architecture also caters for systems with just M and U modes, commonly used in embedded systems, helper cores, and similar use cases. On secure systems not supporting S mode a FW RoT has to share M-mode with an OS. RISC-V does not exclude such implementations - for example, implementations using a certified OS and FW RoT, or using a HW RoT to isolate sensitive code and assets (physical isolation). There is no current mechanism in RISC-V for isolation within M-mode itself other than temporal boundaries

To minimize the TCB of the FW RoT RISC-V recommends that secure systems implement S mode, and de-privilege non-RoT firmware such as an OS.

4.1.2. Isolation model

ID#	Requirement
CAT_NNN	PMP/ePMP, or MTT, MUST be used to isolate M-mode from lower privilege levels.
CAT_NNN	sPMP or MMU MUST be used to isolate user applications or virtual guests



MTT can be sufficient for protecting Root domain in the sense that M-mode can enforce that its own resources are never assigned to another domain. PMP/ePMP still add further protections for M-mode, such as the ability to implement temporal isolation boundaries within M-mode (for example, protect early boot code), or to prevent itself from accessing or executing from memory assigned to lower privilege levels (privilege escalation).

Using sPMP is typically a more static model and can achieve a more deterministic system, for example in automotive or automation use cases.

MMU is typically required for Linux based systems, for example MMI use cases or edge devices.

Either MMU and sPMP can be used both with or without hypervisor extension. For example, the hypervisor extension with sPMP can support static partition hypervisors, commonly used in automotive. And a single stage MMU can be used without hypervisor extension for full Linux support.

4.1.3. Root of Trust

See [reference model](#).

ID#	Requirement
CAT_NNN	A secure system SHOULD implement a HW RoT

4.1.4. Authorized Boot

Multiple models can be used to ensure a secure system can only run authorized software.

See [authorized software](#).

4.1.5. Attestation

Multiple models can be used to prove to a relying party that a secure system is in a trustworthy state.

See [attestable services](#).

4.1.6. Sealing

Multiple models can be used to protect assets if a system is not in a trustworthy state.

See [sealing](#).

4.1.7. Device access control

For the purpose of this specification, a device can be a logical device. A physical device can present one or more logical devices, each with its own (logical) control interface.

Isolation guarantees provided to software also apply to device initiated transaction.

ID#	Requirement
CAT_NNN	IOPMP or IOMTT MUST be used to guarantee that devices assigned to lower privilege levels cannot access resources assigned to M-mode.
CAT_NNN	IOPMP, or IOMTT with IOMMU, MUST be used to enforce access rules for devices assigned to user applications or guests on a virtualized system.

On a non-virtualized system, user devices can be managed by the OS which can enforce access rules for user applications.

On a virtualized system, devices can be virtualized and assigned to guests by the hypervisor configuring MMU and IOMMU translation rules.



IOMTT can also be sufficient for protecting Root devices in the sense that M-mode can enforce that its own resources are never assigned to another domain. Use of IOPMP or similar may still add further protections. For example, a system may require that Root devices cannot access memory assigned to Confidential domain.

4.2. Debug and performance management

See [security lifecycle](#).

See [enhanced RISC-V external debug security](#)

ID#	Requirement
CAT_NNN	External debug SHOULD be enabled separately for M-mode and non M-mode software.
CAT_NNN	External debug MUST only be enabled by HW RoT (M-mode external debug) or by FW RoT (non M-mode external debug).

Enables recoverable external debug of non M-mode software. In turn enabling supply chain separation.

ID#	Requirement
CAT_NNN	Self-hosted debug MAY be used for debug of non M-mode software.

ID#	Requirement
CAT_NNN	Self-hosted debug MUST only be enabled by a higher privileged component.

For example, within normal domain an S-mode OS can enable self-hosted debug for a user application. Only M-mode firmware should enable self-hosted debug for the S-mode OS.

ID#	Requirement
CAT_NNN	FW RoT MAY disable self-hosted debug for all non M-mode software.

For example, disable self-hosted debug in a production system for certification reasons.

ID#	Requirement
CAT_NNN	External debug MUST only be enabled following system reset (part of measuring) of the affected component.
CAT_NNN	Revealing self-hosted debug MUST only be enabled following reboot (part of measuring) of the affected component.
CAT_NNN	Trusted self-hosted debug MAY be enabled at runtime (after measuring) of the affected component, to an application specific governance process.

Guarantees the system remains attestable.

ID#	Requirement
CAT_NNN	Lower privilege software MUST NOT be able to monitor higher privilege software.
CAT_NNN	Software in one domain MUST NOT be able to monitor software in a different domain, without consent.

Prevents using event counters to monitor across application or privilege boundaries. Event counters can be managed by higher privileged software as part of context switching across boundaries.

4.3. Global Platform TEE

4.3.1. Overview

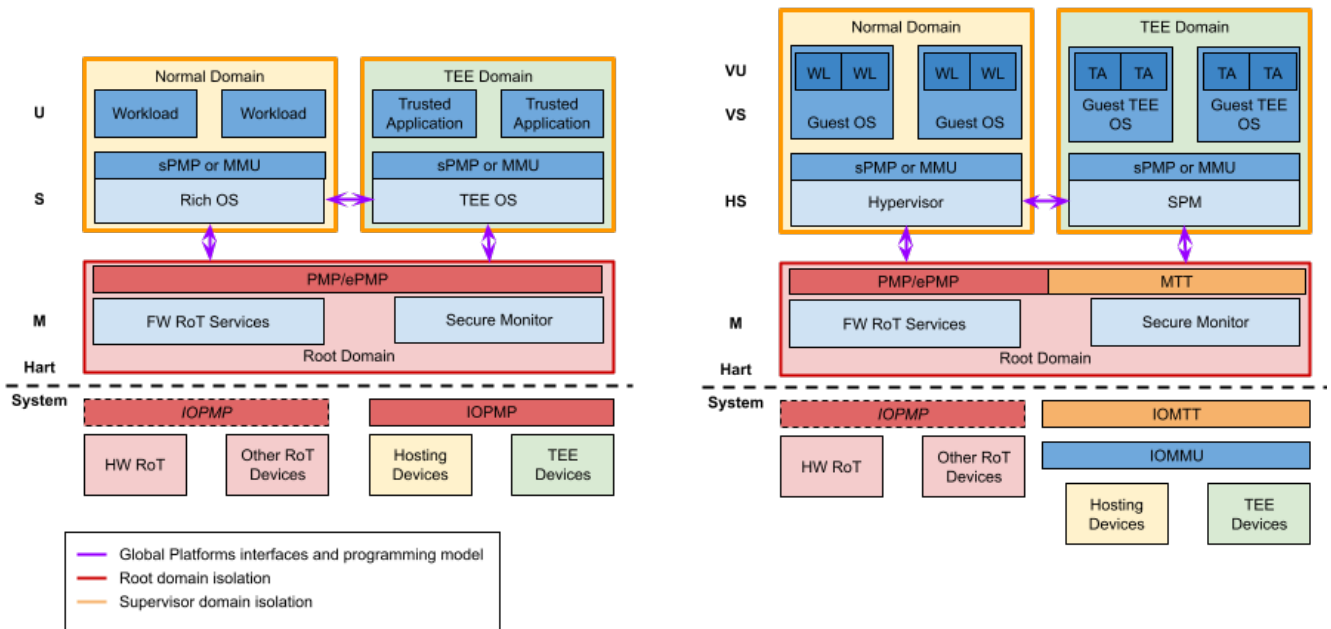


Figure 4: Global platform TEE use cases

Global platform defines technical standards, interface specifications and programming models, open source firmware, and certification programmes for *trusted execution environments (TEE)*.

A TEE is an isolated environment providing security services. TEE services can be available to software on multiple Harts. For example:

- Payment clients
- DRM clients and content protection
- Secure storage
- User identity management
- Attestation services

The TEE model divides software into physically isolated domains:

- Normal domain
Typically hosting a *rich OS* (for example, RTOS or Linux), and user applications.
- TEE domain
Hosts a *TEE OS* (domain security manager) and *trusted applications (TA)*.
- Root domain
Hosts RoT firmware, including a secure monitor.

The TEE OS is primarily responsible for isolation of TA, and for providing root of trust services, within the TEE domain.

The OS in Normal domain typically controls scheduling on the system, across all Harts available to it. To interact with TA services in TEE domain, the OS in Normal domain interacts with a TEE OS through a secure monitor in Root domain.

The secure monitor is responsible for context switching and isolation across domain boundaries, including event management.

For the purpose of this specification, TEE deployment models can be separated as:

- Static partition TEE

A single TEE provides security services to Normal domain. TA are typically installed at boot by RoT FW and TEE OS, though Global Platform does also define protocols for installation of TA at runtime. System configuration and resource allocation can be mostly static, making the system more deterministic.

Use case examples: edge devices and IoT, automation, and automotive.

- Virtualized TEE

On a virtualized system, TEE can also be virtualized. In this case a *secure partition manager* in TEE domain is responsible for isolation of multiple TEE guests (for example, an OEM TEE and separate third party TEE). This model can also support more dynamic resource allocation.

Use case examples: mobile clients, and automotive.

4.3.2. Isolation model

A Global Platform TEE requires the following isolation guarantees:

ID#	Requirement
CAT_NNN	Root domain MAY access resources assigned to any domain, but SHOULD prevent itself from unintended access to resources assigned to a different domain (privilege escalation).
CAT_NNN	No other domains can access resources assigned to Root domain
CAT_NNN	Resources assigned to TEE domain MUST NOT be accessible to Normal domain
CAT_NNN	Resources assigned to Normal domain MUST be accessible to Normal domain (r/w/x), and to TEE domain (r/w) (default sharing rule)
CAT_NNN	Resources assigned to a single TA, or a guest TEE, MUST not be accessible by a different TA, or guest TEE, without consent.

In the standard GP TEE model, each TA is expected to be a self-contained unit providing a specific security service, either to Normal domain or to other TA. All communications are implemented through secure channels managed by the TEE OS or SPM.

Sharing of memory between TA is generally discouraged. But there are mechanisms to do so in specific use cases. For example, sharing media buffers in a secure media path. Such policies are enforced by SPM or TEE OS.

Processes in Normal domain can share memory assigned to Normal domain when interacting with a TA in TEE world (default sharing rule). Such shared memory can be cached when context switching between Normal and TEE domains.

RISC-V hardware enforced isolation mechanisms can be used as follows to meet those guarantees:

ID#	Requirement
CAT_NNN	PMP/ePMP, or MTT, MUST be used to isolate Root domain from other domains.
CAT_NNN	Supervisor domains MUST be used to enforce isolation between Normal and TEE domains.

See [supervisor domains](#).

For static partition TEE, using PMP/ePMP, or PMA, with supervisor domains can be sufficient.

For virtualized TEE, MTT should be used with supervisor domains.



MTT can be sufficient for protecting Root domain in the sense that M-mode can enforce that its own resources are never assigned to another domain. PMP/ePMP still add further protections for M-mode, such as the ability to implement temporal isolation boundaries within M-mode (for example, protect early boot code), or to prevent itself from accessing or executing from memory assigned to lower privilege levels (privilege escalation).

ID#	Requirement
CAT_NNN	For a static partition TEE, sPMP or MMU MUST be used to enforce isolation between TA in TEE domain.

ID#	Requirement
CAT_NNN	For a virtualized TEE, hypervisor extension MUST be supported
CAT_NNN	For a virtualized TEE, MMU MUST be used to enforce isolation between guest TEE, and between TA within a TEE.

4.3.3. Root of Trust

See [reference model](#).

ID#	Requirement
CAT_NNN	A TEE based system SHOULD implement a HW RoT

4.3.4. Authorized boot

See [authorized software](#).

TEE boot is typically based on:

- Measured and verified local boot (direct or indirect)
- Sealing, to protect TEE production assets

The process can involve multiple stages (layered boot).

4.3.5. Attestation

See [attestable services](#).

Static partition TEE attestation is typically based on a direct security platform attestation.

ID#	Requirement
CAT_NNN	A direct security platform attestation MUST cover at least: <ul style="list-style-type: none">• TEE domain• Root domain• Boot state of all trusted subsystems

Virtualized TEE attestation can be layered, for performance or separation of concern. For example:

- A security platform attestation, signed by a RoT, covering trusted subsystems, Root domains, and SPM
- Separate guest TEE attestation(s) signed by SPM

4.3.6. Sealing

See [sealing](#).

In the Global Platform security model, SPM or TEE OS typically provide local trusted storage, key management, and cryptographic services to TA and guest TEE. These services support local sealing of TA or guest TEE assets, and minimize exposure of cryptographic materials.

ID#	Requirement
CAT_NNN	Local sealing for a TA, or a TEE guest, MUST be unique to TEE domain and to a physical instance of a system.
CAT_NNN	Local sealing for a TA, or a TEE guest, SHOULD also be unique to the TEE guest or the TA.
CAT_NNN	Local sealing MAY be layered.

For example:

- TEE domain unique sealing keys derived by a RoT from a hardware unique key
- TA, or guest TEE, unique sealing keys derived by TEE OS or SPM from a TEE domain unique sealing key

4.3.7. Device access control

For the purpose of this specification, a device can be a logical device. A physical device can present one or more logical devices, each with its own (logical) control interface.

The security guarantees also apply to device initiated accesses, for example DMA and interrupts.

ID#	Requirement
CAT_NNN	A static partition TEE MUST use IOPMP to enforce access rules for devices.
CAT_NNN	A virtualized TEE MUST use IOMTT and IOMMU to enforce access rules for devices assigned to Normal or TEE domains, and SHOULD use IOPMP to enforce access rules for Root devices.

For a static partition TEE, domain level granularity can be sufficient as device access within TEE and Normal domains is governed by TEE OS and the rich OS respectively. It can be implemented using IOPMP. Policy can be controlled by boot configuration, by a HW or FW RoT.

For a virtualized TEE, IOMTT enforces supervisor domain level access rules (physical isolation). IOMMU enforces guest and TA level access rules (virtualization), supporting device assignment to a guest TEE or a TA.



IOMTT can also be sufficient for protecting Root devices in the sense that M-mode can enforce that its own resources are never assigned to another domain. Use of IOPMP or similar may still add further protections. For example, a system may require that Root devices cannot be used to access memory assigned to Confidential domain.

4.3.8. System integration

In the case of a Global Platform TEE system a rich OS in Normal domain is free to schedule services, including TEE services, on any Hart available to it. The number and make-up of supervisor domains can be known, and a simple convention can be used for common identification (SDID value, see [supervisor domains](#)) of Normal, TEE, and Root domains across multiple Harts in a system.

System integration in this context involves providing *security attributes* on a system interconnect, tagging all transactions (CPU or system agent initiated) to either Root, Normal, or TEE domains.

Possible use cases include:

- Tweaking cryptographic memory protection (uniqueness)
- Tagging interrupts, debug accesses, or coherent memory accesses
- Device assignment (IOPMP/IOMTT integration), static or dynamic

The attributes can be derived, for example, from SDID and privilege level, from PMA, or from dynamic meta-data during Sv address translation (MTT svpam).

For some use cases security attributes can be extended to reflect finer granularity, for example for cryptographic memory protection with TA granularity.

4.4. Debug and performance management

See [security lifecycle](#).

See [enhanced RISC-V external debug security](#)

ID#	Requirement
CAT_NNN	External debug MUST be enabled separately for Root domain.
CAT_NNN	External debug MUST be enabled separately for each supervisor domain.
CAT_NNN	External debug MUST only be enabled by a HW RoT (Root domain external debug) or by Root domain (supervisor domain external debug).

Enables recoverable external debug of a supervisor domain separately from other supervisor domains, and Root domain. In turn enabling supply chain separation.

ID#	Requirement
CAT_NNN	Self-hosted debug MAY be used for debug within a supervisor domain.
CAT_NNN	Self-hosted debug MUST only be enabled by a higher privileged component.

For example, within normal domain an S-mode or VS-mode OS can enable self-hosted debug for a user application. Or an HS-mode hypervisor can enable self-hosted debug for a VS-mode guest. Only Root domain should enable self-hosted debug for an S-mode OS or an HS mode hypervisor.

Within TEE domain a TEE OS can enable self-hosted debug for a TA. An SPM can enable self-hosted debug for guest TEE. Only Root domain should enable self-hosted debug of SPM (virtualized) or TEE OS (non-virtualized).

ID#	Requirement
CAT_NNN	Root domain MAY disable self-hosted debug for a whole domain.

For example, for all of TEE domain on a production system, for certification reasons.

ID#	Requirement
CAT_NNN	External debug MUST only be enabled following system reset (part of measuring) of the affected component.
CAT_NNN	Revealing self-hosted debug MUST only be enabled following reboot (part of measuring) of the affected component.
CAT_NNN	Trusted self-hosted debug MAY be enabled at runtime (after measuring) of the affected component, to an application specific governance process.

Guarantees the system remains attestable.

ID#	Requirement
CAT_NNN	Lower privilege software MUST NOT be able to monitor higher privilege software.
CAT_NNN	Software in one domain MUST NOT be able to monitor software in a different domain, without consent.

Prevents using event counters to monitor across guest/application, privilege and supervisor domain boundaries. Event counters can be managed by higher privileged software as part of context

switching across boundaries.

4.5. Confidential computing on RISC-V (CoVE)

4.5.1. Overview

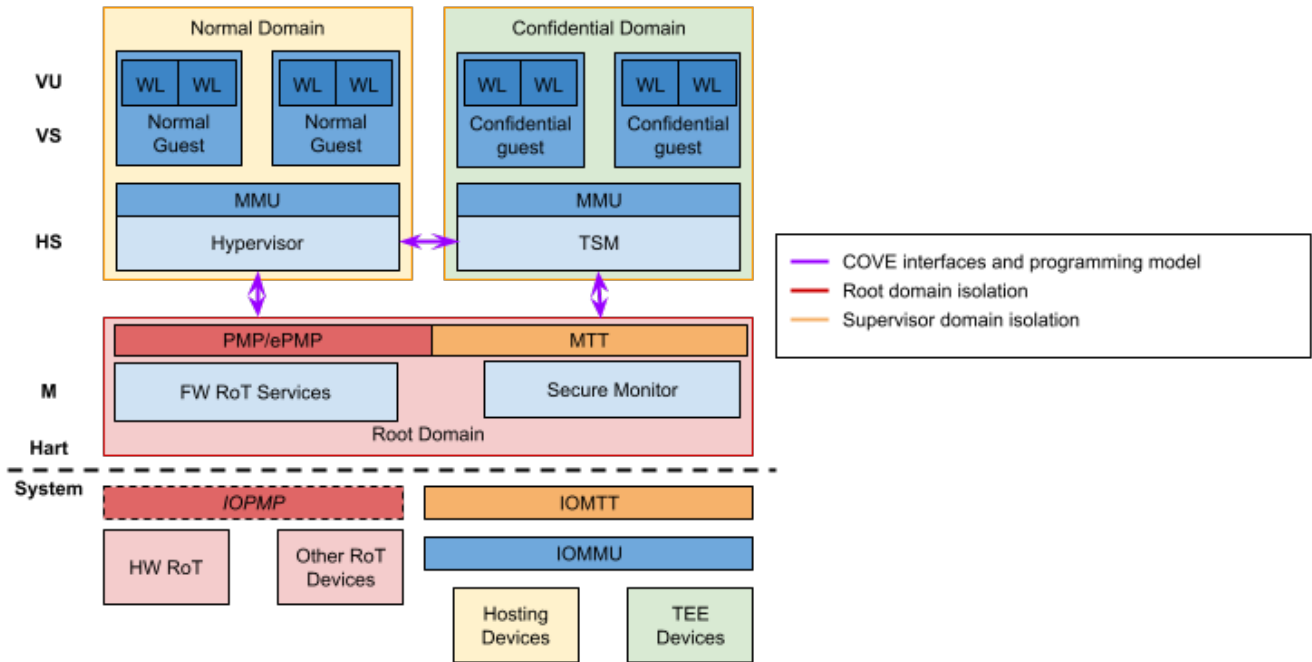


Figure 5: Confidential compute use case

In hosting environments, tenant workloads rely on isolation primitives that are managed by host privileged software. This can lead to a large TCB for tenants which could include, for example, a hypervisor, orchestration services, and host management services. It could also include other tenants exploiting vulnerabilities in complex hosting software.

Confidential compute aims to achieve a minimal and certifiable TCB for *confidential workloads*.

CoVE (Confidential VM Extensions) [specification](#) defines a confidential compute platform for RISC-V systems, including interfaces and programming models, covering lifecycle management, attestation, resource management and devices assignment, for confidential workloads. It is based on principles defined by [Confidential Computing Consortium](#). Reference firmware for CoVE is being developed as part of the [RISC-V Software Ecosystem](#) project.

CoVE is primarily aimed at cloud hosting of confidential workloads. But the underlying isolation model could potentially be used in other use cases, such as some mobile clients or edge devices.

CoVE divides software into physically isolated domains:

- Normal domain
Typically hosting a hypervisor, and Normal guests and services.
- Confidential domain
Hosts a *TSM* (domain security manager) and confidential guests.
- Root domain

Hosts RoT firmware, including a secure monitor.

The TSM is primarily responsible for isolation of confidential workloads, and for providing RoT services, within the Confidential domain.

A hypervisor in Normal domain typically controls scheduling and resource assignment on the system across all Harts available to it, including for confidential workloads. It interacts with the TSM through the secure monitor in Root domain to manage confidential workloads.

The secure monitor is responsible for context switching and isolation across domain boundaries, including event management.

4.5.2. Isolation model

Confidential workloads are provided the following isolation guarantees:

ID#	Requirement
CAT_NNN	Root domain MAY access resources assigned to any domain, but SHOULD prevent itself from unintended access to resources assigned to a different domain (privilege escalation).
CAT_NNN	Resources assigned to Root domain MUST be private to Root domain
CAT_NNN	Resources assigned only to Confidential domain MUST not be accessible by Normal domain
CAT_NNN	Resources assigned only to Normal domain MUST not be accessible by Confidential domain
CAT_NNN	Resources MAY be assigned to both Normal and Confidential domains (sharing by consent).
CAT_NNN	Resources assigned to a single confidential workload MUST NOT be accessible by any other confidential workload
CAT_NNN	Resources MAY be assigned to multiple confidential workloads (sharing by consent)

RISC-V hardware enforced isolation mechanisms can be used as follows to meet those guarantees:

ID#	Requirement
CAT_NNN	PMP/ePMP or MTT MUST be used to isolate Root domain from other domains.
CAT_NNN	Supervisor domains MUST be used to enforce isolation between Normal and Confidential domains.

See [supervisor domains](#).



MTT can be sufficient for protecting Root domain in the sense that M-mode can enforce that its own resources are never assigned to another domain. PMP/ePMP still add further protections for M-mode, such as the ability to implement temporal isolation boundaries within M-mode (for example, protect early boot code), or to

prevent itself from accessing or executing from memory assigned to lower privilege levels (privilege escalation).

ID#	Requirement
CAT_NNN	Hypervisor extension MUST be supported
CAT_NNN	MMU MUST be used to enforce isolation between Confidential guests within Confidential domain.

4.5.3. Root of trust

See [reference model](#).

ID#	Requirement
CAT_NNN	A CoVE system MUST implement a HW RoT

4.5.4. Authorized Boot

See [authorized software](#).

ID#	Requirement
CAT_NNN	Confidential guests MUST not boot until at least the security platform has been verified: <ul style="list-style-type: none">• TSM in Confidential domain• Root domain• Boot state of all trusted subsystems

Boot in a cloud hosting context is typically based on:

- Measured boot of a hosting platform, including Root domain and TSM
- Platform attestation and security provisioning (unsealing) by a remote provisioning system
- Launch and measurement of confidential workloads, only once the system has been unsealed

A *trusted platform module* (TPM) can be used to measure the security platform.

Measuring confidential guests can be done by TSM in Confidential domain.

The process can involve multiple stages (layered boot).

4.5.5. Attestation

See [attestable services](#).

Attestation of confidential workloads is typically layered, for performance and separation of concern:

- A security platform attestation, signed by a hardware root of trust
- A confidential workload attestation, signed by TSM

ID#	Requirement
CAT_NNN	<p>A security platform attestation MUST cover at least:</p> <ul style="list-style-type: none"> • HW RoT • TSM • Root domain • Boot state of all trusted subsystems

4.5.6. Sealing

See [sealing](#).

Sealing of confidential workloads is typically based on remote sealing, unsealing assets for a confidential workload following successful attestation by a remote provisioning system. This enables use cases such as:

- Shared assets across multiple instances of a confidential workload (scale or redundancy)
- Unsealing different sets of assets for different users of a service

TSM itself is typically stateless across reset and does not require any sealed assets of its own.

4.5.7. Device access control

For the purpose of this specification, a device can be a logical device. A physical device can present more than one logical devices, each with its own (logical) control interface.

The security guarantees also apply to device initiated accesses, for example DMA and interrupts.

ID#	Requirement
CAT_NNN	IOMTT and IOMMU MUST be used to enforce access rules for devices assigned to Normal or Confidential domains.
CAT_NNN	IOPMP SHOULD be used to enforce access rules for Root devices.
CAT_NNN	IOPMP and IOMTT configurations MUST only be directly accessible by Root domain.

IOMTT enforces supervisor domain level access rules (physical isolation). IOMMU enforces guest and TA level access rules (virtualization), supporting device assignment to a Confidential guest.



IOMTT can also be sufficient for protecting Root devices in the sense that M-mode can enforce that its own resources are never assigned to another domain. Use of IOPMP or similar still adds further protections. For example, a system may require that Root devices cannot be used to access memory assigned to Confidential

domain.

4.5.8. System integration

In the case of a confidential compute system, hypervisor in Normal domain typically controls scheduling and resource assignment on the system across all Harts available to it. The number and make-up of supervisor domains can be known, and a simple convention can be used for common identification of Normal, Confidential, and Root domains across multiple Harts in a system.

System integration in this context involves providing *security attributes* on the interconnect, tagging all transactions (CPU or system agent initiated) to either Root, Normal, or TEE domains.

Possible use cases include:

- Tweaking cryptographic memory protection (uniqueness)
- Tagging interrupts, debug accesses, or coherent memory accesses
- Device assignment (IOPMP/IOMTT integration), static or dynamic

The attributes can be derived, for example, from dynamic meta-data during Sv address translation (MTT Svpm).

For some use cases security attributes can be extended to reflect finer granularity, for example for cryptographic memory protection with confidential workload granularity.

4.5.9. Trusted device assignment

The goal of confidential compute is to provide a minimum TCB for a confidential service, and CPU isolation mechanisms discussed so far does that on a Hart.

But most confidential services also make use of devices, both on-chip and external. [Device virtualization](#) can guarantee exclusivity for devices assigned to a confidential workload - TSM can guarantee that a device assigned to a confidential workload cannot be accessed by:

- Any other confidential workload
- Any software in Normal domain

But the confidential workload still has to trust all intermediaries between the workload and the device, both physical and software. For example:

- Drivers
- Physical interconnects and device hardware interfaces

Secure access to devices is important in a number of use cases where a device performs work on assets owned by a confidential workload, such as accelerators.

The *TEE device interface security protocol (TDISP)* defined by PCIe provides a security architecture and protocols allowing a confidential workload to securely attest, manage and exchange data with a trusted device.

CoVE defines RISC-V support for TDISP. See:

pcisig.com/specifications/ github.com/riscv-non-isa/riscv-ap-tee-io

4.5.10. Debug and performance management

See [security lifecycle](#).

See [enhanced RISC-V external debug security](#)

ID#	Requirement
CAT_NNN	External debug MUST be enabled separately for Root domain.
CAT_NNN	External debug MUST be enabled separately for each supervisor domain.
CAT_NNN	External debug MUST only be enabled by a HW RoT (Root domain external debug) or by Root domain (supervisor domain external debug).

Enables recoverable external debug of a supervisor domain separately from other supervisor domains, and Root domain. In turn enabling supply chain separation.

ID#	Requirement
CAT_NNN	Self-hosted debug MAY be used for debug within a supervisor domain.
CAT_NNN	Self-hosted debug MUST only be enabled by a higher privileged component.

For example, within normal domain an HS-mode hypervisor can enable self-hosted debug for a VS-mode guest. Only Root domain should enable self-hosted debug for the HS mode hypervisor.

Within Confidential domain the TSM can enable self-hosted debug for a confidential guest. Only Root domain should enable self-hosted debug of TSM.

ID#	Requirement
CAT_NNN	External debug MUST only be enabled following system reset (part of measuring) of the affected component.
CAT_NNN	Revealing self-hosted debug MUST only be enabled following reboot (part of measuring) of the affected component.
CAT_NNN	Trusted self-hosted debug MAY be enabled at runtime (after measuring) of the affected component, to an application specific governance process.

Guarantees the system remains attestable.

ID#	Requirement
CAT_NNN	Lower privilege software MUST NOT be able to measure higher privilege software.
CAT_NNN	Software in one domain MUST NOT be able to measure software in a different domain, without consent.

Prevents using event counters to measure across guest/application, privilege and supervisor

domain boundaries.

Event counters can be managed by higher privileged software as part of context switching across boundaries.

4.5.11. Platform QoS

See [platform quality of service](#).

ID#	Requirement
CAT_NNN	Lower privilege software MUST NOT be able to measure higher privilege software.
CAT_NNN	Software in one domain MUST NOT be able to measure software in a different domain, without consent.

Event counters can be managed by higher privileged software as part of context switching across boundaries.

Chapter 5. Cryptography

<use models enabled by crypto ISA extensions>

<crypto guidelines>

Appendix A: References

1. <https://www.intel.com/content/www/us/en/newsroom/opinion/zero-trust-approach-architecting-silicon.html>
2. <https://www.forrester.com/blogs/tag/zero-trust/>
3. <https://docs.microsoft.com/en-us/security/zero-trust/>
4. <https://github.com/riscv/riscv-crypto/releases>
5. <https://github.com/riscv/riscv-platform-specs/blob/main/riscv-platform-spec.adoc>
6. https://www.commoncriteriaportal.org/files/ppfiles/pp0084b_pdf.pdf
7. https://docs.opentitan.org/doc/security/specs/device_life_cycle/
8. <https://nvlpubs.nist.gov/nistpubs/ir/2021/NIST.IR.8320-draft.pdf>
9. <https://nvlpubs.nist.gov/nistpubs/SpecialPublications/NIST.SP.800-193.pdf>
10. <https://www.rambus.com/security/root-of-trust/rt-630/>
11. <https://docs.opentitan.org/doc/security/specs/>
12. <https://trustedcomputinggroup.org/work-groups/dice-architectures/>
13. <https://ieeexplore.ieee.org/iel7/8168766/8203442/08203496.pdf>
14. <https://dl.acm.org/doi/10.1145/168619.168635>
15. <https://dl.acm.org/doi/abs/10.1145/3342195.3387532>
16. <https://github.com/riscv/riscv-debug-spec/blob/master/riscv-debug-stable.pdf>
17. https://csrc.nist.gov/csrc/media/events/non-invasive-attack-testing-workshop/documents/08_goodwill.pdf
18. <https://www.iso.org/standard/60612.html>
19. <https://ieeexplore.ieee.org/document/6176671>
20. <https://tches.iacr.org/index.php/TCHES/article/view/8988>
21. <https://ieeexplore.ieee.org/abstract/document/1401864>
22. <https://www.electronicsspecifier.com/products/design-automation/increasingly-connected-world-needs-greater-security>
23. <https://www.samsungknox.com/es-419/blog/knox-e-fota-and-sequential-updates>
24. <https://docs.microsoft.com/en-us/windows/security/threat-protection/intelligence/supply-chain-malware>
25. <https://dl.acm.org/doi/10.1145/3466752.3480068>
26. <https://arxiv.org/abs/2111.01421>
27. <https://www.nap.edu/catalog/24676/foundational-cybersecurity-research-improving-science-engineering-and-institutions>
28. <https://trustedcomputinggroup.org/work-groups/dice-architectures/>

Bibliography

- [1] The RISC-V Instruction Set Manual Volume II: Privileged Architecture Document Version 20211203 ([link](#))